

DSC680_Project1Final_LincolnBrown

December 22, 2024

0.1 Import the libraries

```
[190]: import category_encoders as ce
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.metrics import roc_curve, auc, roc_auc_score
from sklearn.model_selection import train_test_split
from sklearn.utils.class_weight import compute_class_weight
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Dropout
from tensorflow.keras.preprocessing.sequence import pad_sequences
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix
from sklearn.preprocessing import StandardScaler
import tensorflow as tf
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.layers import LSTM, Dense, Dropout
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.optimizers.legacy import Adam
import time
```

0.2 Import the Dataset

```
[2]: cc_f = "creditcard_dataset/credit_card_transactions-ibm_v2.csv"
user_f = "creditcard_dataset/sd254_users.csv"
cc_df = pd.read_csv(cc_f)
user_df = pd.read_csv(user_f)
```

```
[3]: cc_df
```

```
[3]:
```

	User	Card	Year	Month	Day	Time	Amount	Use Chip	\
0	0	0	2002	9	1	06:21	\$134.09	Swipe Transaction	
1	0	0	2002	9	1	06:42	\$38.48	Swipe Transaction	
2	0	0	2002	9	2	06:22	\$120.34	Swipe Transaction	
3	0	0	2002	9	2	17:45	\$128.95	Swipe Transaction	

4	0	0	2002	9	3	06:23	\$104.71	Swipe Transaction
...
24386895	1999	1	2020	2	27	22:23	\$-54.00	Chip Transaction
24386896	1999	1	2020	2	27	22:24	\$54.00	Chip Transaction
24386897	1999	1	2020	2	28	07:43	\$59.15	Chip Transaction
24386898	1999	1	2020	2	28	20:10	\$43.12	Chip Transaction
24386899	1999	1	2020	2	28	23:10	\$45.13	Chip Transaction

	Merchant Name	Merchant City	Merchant State	Zip	MCC	\
0	3527213246127876953	La Verne	CA	91750.0	5300	
1	-727612092139916043	Monterey Park	CA	91754.0	5411	
2	-727612092139916043	Monterey Park	CA	91754.0	5411	
3	3414527459579106770	Monterey Park	CA	91754.0	5651	
4	5817218446178736267	La Verne	CA	91750.0	5912	
...	
24386895	-5162038175624867091	Merrimack	NH	3054.0	5541	
24386896	-5162038175624867091	Merrimack	NH	3054.0	5541	
24386897	2500998799892805156	Merrimack	NH	3054.0	4121	
24386898	2500998799892805156	Merrimack	NH	3054.0	4121	
24386899	4751695835751691036	Merrimack	NH	3054.0	5814	

	Errors?	Is Fraud?
0	NaN	No
1	NaN	No
2	NaN	No
3	NaN	No
4	NaN	No
...
24386895	NaN	No
24386896	NaN	No
24386897	NaN	No
24386898	NaN	No
24386899	NaN	No

[24386900 rows x 15 columns]

```
[4]: user_df
```

```
[4]:
```

	Person	Current Age	Retirement Age	Birth Year	Birth Month	\
0	Hazel Robinson	53	66	1966	11	
1	Sasha Sadr	53	68	1966	12	
2	Saanvi Lee	81	67	1938	11	
3	Everlee Clark	63	63	1957	1	
4	Kyle Peterson	43	70	1976	9	
...	
1995	Jose Faraday	32	70	1987	7	
1996	Ximena Richardson	62	65	1957	11	

1997	Annika Russell	47	67	1973	1
1998	Juelz Roman	66	60	1954	2
1999	Kenia Harris	21	60	1998	11

	Gender	Address	Apartment	City	State	\
0	Female	462 Rose Lane	NaN	La Verne	CA	
1	Female	3606 Federal Boulevard	NaN	Little Neck	NY	
2	Female	766 Third Drive	NaN	West Covina	CA	
3	Female	3 Madison Street	NaN	New York	NY	
4	Male	9620 Valley Stream Drive	NaN	San Francisco	CA	
...	
1995	Male	6577 Lexington Lane	9.0	Freeport	NY	
1996	Female	2 Elm Drive	955.0	Independence	KY	
1997	Female	276 Fifth Boulevard	NaN	Elizabeth	NJ	
1998	Male	259 Valley Boulevard	NaN	Camp Hill	PA	
1999	Female	472 Ocean View Street	NaN	Merrimack	NH	

	Zipcode	Latitude	Longitude	Per Capita Income	Zipcode	\
0	91750	34.15	-117.76		\$29278	
1	11363	40.76	-73.74		\$37891	
2	91792	34.02	-117.89		\$22681	
3	10069	40.71	-73.99		\$163145	
4	94117	37.76	-122.44		\$53797	
...	
1995	11520	40.65	-73.58		\$23550	
1996	41051	38.95	-84.54		\$24218	
1997	7201	40.66	-74.19		\$15175	
1998	17011	40.24	-76.92		\$25336	
1999	3054	42.86	-71.48		\$32325	

	Yearly Income - Person	Total Debt	FICO Score	Num Credit Cards
0	\$59696	\$127613	787	5
1	\$77254	\$191349	701	5
2	\$33483	\$196	698	5
3	\$249925	\$202328	722	4
4	\$109687	\$183855	675	1
...
1995	\$48010	\$87837	703	3
1996	\$49378	\$104480	740	4
1997	\$30942	\$71066	779	3
1998	\$54654	\$27241	618	1
1999	\$65909	\$181261	673	2

[2000 rows x 18 columns]

0.3 Check/Clean for missing Values

```
[5]: cc_df.loc[cc_df['Merchant State'].isna()]
```

```
[5]:
```

	User	Card	Year	Month	Day	Time	Amount	Use Chip	\
11	0	0	2002	9	5	20:41	\$53.91	Online Transaction	
24	0	0	2002	9	9	20:02	\$144.90	Online Transaction	
85	0	0	2002	9	30	06:21	\$127.32	Online Transaction	
99	0	0	2002	10	6	06:14	\$139.39	Online Transaction	
106	0	0	2002	10	9	08:16	\$53.09	Online Transaction	
...	
24386877	1999	1	2020	2	24	20:04	\$55.79	Online Transaction	
24386879	1999	1	2020	2	25	07:06	\$43.08	Online Transaction	
24386880	1999	1	2020	2	25	07:34	\$43.76	Online Transaction	
24386884	1999	1	2020	2	26	07:43	\$45.18	Online Transaction	
24386889	1999	1	2020	2	27	07:47	\$47.18	Online Transaction	


	Merchant Name	Merchant City	Merchant State	Zip	MCC	Errors?	\
11	-9092677072201095172	ONLINE	NaN	NaN	4900	NaN	
24	-8338381919281017248	ONLINE	NaN	NaN	4899	NaN	
85	-7421093378627544099	ONLINE	NaN	NaN	5311	NaN	
99	-7421093378627544099	ONLINE	NaN	NaN	5311	NaN	
106	-4956618006720593695	ONLINE	NaN	NaN	5193	NaN	
...	
24386877	-6160036380778658394	ONLINE	NaN	NaN	4121	NaN	
24386879	-6160036380778658394	ONLINE	NaN	NaN	4121	NaN	
24386880	-6160036380778658394	ONLINE	NaN	NaN	4121	NaN	
24386884	-6160036380778658394	ONLINE	NaN	NaN	4121	NaN	
24386889	-5841929396161652653	ONLINE	NaN	NaN	4121	NaN	

	Is Fraud?
11	No
24	No
85	No
99	No
106	No
...	...
24386877	No
24386879	No
24386880	No
24386884	No
24386889	No

```
[2720821 rows x 15 columns]
```

It looks like all of the missing values for Merchant State are due to Online purchases, we will resolve this by labeling these missing values as ONLINE, similar to what is already found in Merchant City

```
[6]: percent_missing=(cc_df.isnull().sum()*100/cc_df.shape[0]).
      ↪sort_values(ascending=True)
plt.title("Missing Value Analysis")
plt.xlabel("Features")
plt.ylabel("% of missing values")
plt.bar(percent_missing.sort_values(ascending=False).index,percent_missing.
      ↪sort_values(ascending=False),color=(0.1, 0.1, 0.1, 0.1),edgecolor='blue')
plt.xticks(rotation=90)
plt.show()
```



```
[7]: cc_df
```

```
[7]:      User  Card  Year  Month  Day  Time  Amount  Use Chip \
0         0    0  2002     9    1  06:21  $134.09  Swipe Transaction
```

1	0	0	2002	9	1	06:42	\$38.48	Swipe Transaction
2	0	0	2002	9	2	06:22	\$120.34	Swipe Transaction
3	0	0	2002	9	2	17:45	\$128.95	Swipe Transaction
4	0	0	2002	9	3	06:23	\$104.71	Swipe Transaction
...
24386895	1999	1	2020	2	27	22:23	\$-54.00	Chip Transaction
24386896	1999	1	2020	2	27	22:24	\$54.00	Chip Transaction
24386897	1999	1	2020	2	28	07:43	\$59.15	Chip Transaction
24386898	1999	1	2020	2	28	20:10	\$43.12	Chip Transaction
24386899	1999	1	2020	2	28	23:10	\$45.13	Chip Transaction

	Merchant Name	Merchant City	Merchant State	Zip	MCC	\
0	3527213246127876953	La Verne	CA	91750.0	5300	
1	-727612092139916043	Monterey Park	CA	91754.0	5411	
2	-727612092139916043	Monterey Park	CA	91754.0	5411	
3	3414527459579106770	Monterey Park	CA	91754.0	5651	
4	5817218446178736267	La Verne	CA	91750.0	5912	
...
24386895	-5162038175624867091	Merrimack	NH	3054.0	5541	
24386896	-5162038175624867091	Merrimack	NH	3054.0	5541	
24386897	2500998799892805156	Merrimack	NH	3054.0	4121	
24386898	2500998799892805156	Merrimack	NH	3054.0	4121	
24386899	4751695835751691036	Merrimack	NH	3054.0	5814	

	Errors?	Is Fraud?
0	NaN	No
1	NaN	No
2	NaN	No
3	NaN	No
4	NaN	No
...
24386895	NaN	No
24386896	NaN	No
24386897	NaN	No
24386898	NaN	No
24386899	NaN	No

[24386900 rows x 15 columns]

```
[8]: cc_df['Amount'] = cc_df['Amount'].str.replace("\\$", "", regex=True)
cc_df['Amount'] = pd.to_numeric(cc_df['Amount'], errors='coerce')
print(cc_df['Amount'].head(10))
```

0	134.09
1	38.48
2	120.34
3	128.95
4	104.71

```

5      86.19
6      93.84
7     123.50
8      61.72
9      57.10
Name: Amount, dtype: float64

```

```
[9]: cc_df['Errors?'].unique()
```

```
[9]: array([nan, 'Technical Glitch', 'Insufficient Balance', 'Bad PIN',
        'Bad PIN,Insufficient Balance', 'Bad Expiration',
        'Bad PIN,Technical Glitch', 'Bad Card Number', 'Bad CVV',
        'Bad Zipcode', 'Insufficient Balance,Technical Glitch',
        'Bad Card Number,Insufficient Balance', 'Bad Card Number,Bad CVV',
        'Bad CVV,Insufficient Balance', 'Bad Card Number,Bad Expiration',
        'Bad Expiration,Bad CVV', 'Bad Expiration,Insufficient Balance',
        'Bad Expiration,Technical Glitch',
        'Bad Card Number,Bad Expiration,Technical Glitch',
        'Bad CVV,Technical Glitch', 'Bad Card Number,Technical Glitch',
        'Bad Zipcode,Insufficient Balance', 'Bad Zipcode,Technical Glitch',
        'Bad Card Number,Bad Expiration,Insufficient Balance'],
        dtype=object)
```

0.4 Clean the dataset

```
[10]: cc_df['Zip'] = cc_df['Zip'].astype(str)
cc_df.loc[cc_df['Merchant City'] == 'ONLINE', ['Merchant State', 'Zip']] =
↳ 'ONLINE'
```

```
[11]: cc_df.loc[cc_df['Zip'].isna(), 'Zip'] = 'Foreign'
```

```
[12]: cc_df.isna().any()
```

```
[12]: User           False
Card              False
Year             False
Month            False
Day              False
Time             False
Amount           False
Use Chip         False
Merchant Name     False
Merchant City     False
Merchant State    False
Zip              False
MCC              False
Errors?          True
Is Fraud?        False
```

dtype: bool

```
[13]: # Get a unique list of "States"  
states = cc_df['Merchant State'].unique()
```

```
[14]: # Check the US States abbreviations  
us_states = [state for state in states if len(str(state)) == 2]
```

```
[15]: us_states
```

```
[15]: ['CA',  
      'NE',  
      'IL',  
      'MO',  
      'IA',  
      'TX',  
      'NJ',  
      'NV',  
      'NY',  
      'AZ',  
      'UT',  
      'FL',  
      'MI',  
      'WA',  
      'OH',  
      'NM',  
      'SC',  
      'AK',  
      'PA',  
      'VA',  
      'HI',  
      'CT',  
      'MA',  
      'MN',  
      'CO',  
      'GA',  
      'AR',  
      'OR',  
      'WI',  
      'NC',  
      'WV',  
      'ME',  
      'NH',  
      'VT',  
      'MD',  
      'AL',  
      'KY',
```



```
'TN',
'MS',
'KS',
'ND',
'DC',
'MT',
'OK',
'WY',
'ID',
'RI',
'IN',
'LA',
'DE',
'SD',
'AA']
```

```
[16]: len(us_states)
```

```
[16]: 52
```

It looks like there are 52 state abbreviations, including DC for District of Columbia and AA for Armed Forces of America

Let's map the days of the month to day of the week and separate the time column into hour and minute

```
[17]: # Get the day of the month

cc_df['Date'] = pd.to_datetime(cc_df[['Year', 'Month', 'Day']])

# Extract day of the week and map it to its name
days = {0:'Mon', 1:'Tue', 2:'Wed', 3:'Thu', 4:'Fri', 5:'Sat', 6:'Sun'}
cc_df['Day of Week'] = cc_df['Date'].dt.dayofweek.map(days)
```

```
[18]: # Get the hour and minute from Time
cc_df['Hour'] = pd.to_numeric(cc_df['Time'].str[0:2])
cc_df['Minute'] = pd.to_numeric(cc_df['Time'].str[3:5])
```

```
[19]: cc_df
```

```
[19]:
```

	User	Card	Year	Month	Day	Time	Amount	Use Chip	\
0	0	0	2002	9	1	06:21	134.09	Swipe Transaction	
1	0	0	2002	9	1	06:42	38.48	Swipe Transaction	
2	0	0	2002	9	2	06:22	120.34	Swipe Transaction	
3	0	0	2002	9	2	17:45	128.95	Swipe Transaction	
4	0	0	2002	9	3	06:23	104.71	Swipe Transaction	
...
24386895	1999	1	2020	2	27	22:23	-54.00	Chip Transaction	

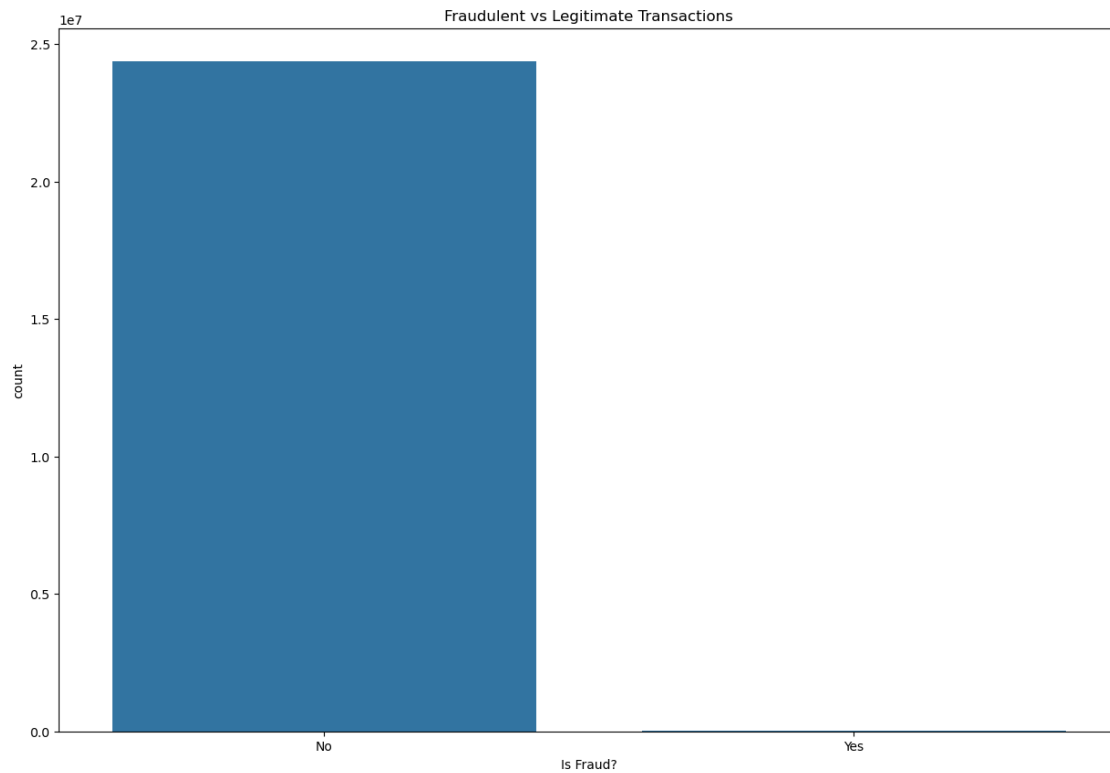
24386896	1999	1	2020	2	27	22:24	54.00	Chip Transaction
24386897	1999	1	2020	2	28	07:43	59.15	Chip Transaction
24386898	1999	1	2020	2	28	20:10	43.12	Chip Transaction
24386899	1999	1	2020	2	28	23:10	45.13	Chip Transaction

	Merchant Name	Merchant City	Merchant State	Zip	MCC	\
0	3527213246127876953	La Verne	CA	91750.0	5300	
1	-727612092139916043	Monterey Park	CA	91754.0	5411	
2	-727612092139916043	Monterey Park	CA	91754.0	5411	
3	3414527459579106770	Monterey Park	CA	91754.0	5651	
4	5817218446178736267	La Verne	CA	91750.0	5912	
...	
24386895	-5162038175624867091	Merrimack	NH	3054.0	5541	
24386896	-5162038175624867091	Merrimack	NH	3054.0	5541	
24386897	2500998799892805156	Merrimack	NH	3054.0	4121	
24386898	2500998799892805156	Merrimack	NH	3054.0	4121	
24386899	4751695835751691036	Merrimack	NH	3054.0	5814	

	Errors?	Is Fraud?	Date	Day of Week	Hour	Minute
0	NaN	No	2002-09-01	Sun	6	21
1	NaN	No	2002-09-01	Sun	6	42
2	NaN	No	2002-09-02	Mon	6	22
3	NaN	No	2002-09-02	Mon	17	45
4	NaN	No	2002-09-03	Tue	6	23
...
24386895	NaN	No	2020-02-27	Thu	22	23
24386896	NaN	No	2020-02-27	Thu	22	24
24386897	NaN	No	2020-02-28	Fri	7	43
24386898	NaN	No	2020-02-28	Fri	20	10
24386899	NaN	No	2020-02-28	Fri	23	10

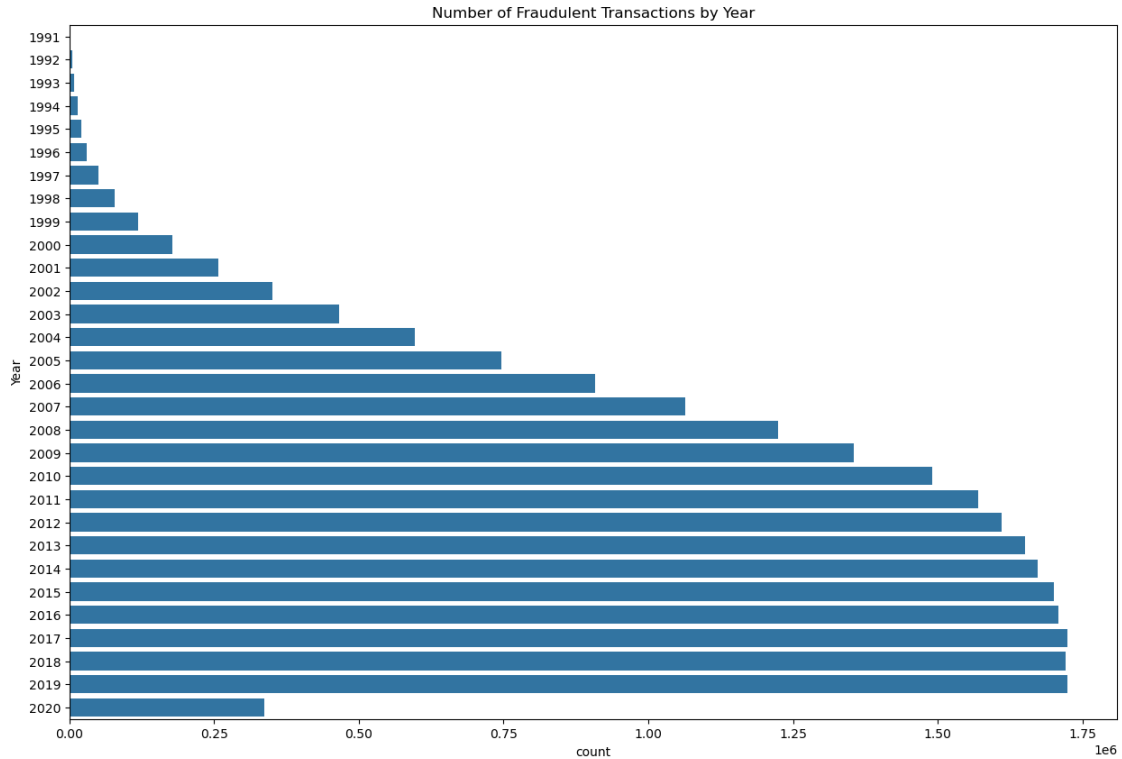
[24386900 rows x 19 columns]

```
[20]: # Fraudulent vs legitimate transactions
plt.figure(figsize=(15, 10))
sns.countplot(data=cc_df, x='Is Fraud?')
plt.title('Fraudulent vs Legitimate Transactions')
plt.show()
```



As we can see, this dataset is heavily imbalanced, with the number of legitimate transactions far exceeding fraudulent.

```
[21]: plt.figure(figsize=(15, 10))
sns.countplot(data=cc_df, y='Year')
plt.title('Number of Fraudulent Transactions by Year')
plt.show()
```



It looks like the number of fraudulent transactions increases year over year. I want to confirm that we do not have full yearly data for 2020.

```
[22]: cc_df['Date'].max()
```

```
[22]: Timestamp('2020-02-28 00:00:00')
```

The latest date we have for 2020 is February 28th, indicating that we do not have a full year's worth of data for 2020.

0.5 Fraudulent Transactions Analysis

Let's take a look at the amounts of the fraudulent transactions.

```
[23]: cc_df['Is Fraud?'] = cc_df['Is Fraud?'].map({'No': 0, 'Yes': 1})
cc_df
```

```
[23]:
```

	User	Card	Year	Month	Day	Time	Amount	Use Chip	\
0	0	0	2002	9	1	06:21	134.09	Swipe Transaction	
1	0	0	2002	9	1	06:42	38.48	Swipe Transaction	
2	0	0	2002	9	2	06:22	120.34	Swipe Transaction	
3	0	0	2002	9	2	17:45	128.95	Swipe Transaction	
4	0	0	2002	9	3	06:23	104.71	Swipe Transaction	
...	

24386895	1999	1	2020	2	27	22:23	-54.00	Chip Transaction
24386896	1999	1	2020	2	27	22:24	54.00	Chip Transaction
24386897	1999	1	2020	2	28	07:43	59.15	Chip Transaction
24386898	1999	1	2020	2	28	20:10	43.12	Chip Transaction
24386899	1999	1	2020	2	28	23:10	45.13	Chip Transaction

	Merchant Name	Merchant City	Merchant State	Zip	MCC	\
0	3527213246127876953	La Verne	CA	91750.0	5300	
1	-727612092139916043	Monterey Park	CA	91754.0	5411	
2	-727612092139916043	Monterey Park	CA	91754.0	5411	
3	3414527459579106770	Monterey Park	CA	91754.0	5651	
4	5817218446178736267	La Verne	CA	91750.0	5912	
...	
24386895	-5162038175624867091	Merrimack	NH	3054.0	5541	
24386896	-5162038175624867091	Merrimack	NH	3054.0	5541	
24386897	2500998799892805156	Merrimack	NH	3054.0	4121	
24386898	2500998799892805156	Merrimack	NH	3054.0	4121	
24386899	4751695835751691036	Merrimack	NH	3054.0	5814	

	Errors?	Is Fraud?	Date	Day	of Week	Hour	Minute
0	NaN	0	2002-09-01		Sun	6	21
1	NaN	0	2002-09-01		Sun	6	42
2	NaN	0	2002-09-02		Mon	6	22
3	NaN	0	2002-09-02		Mon	17	45
4	NaN	0	2002-09-03		Tue	6	23
...
24386895	NaN	0	2020-02-27		Thu	22	23
24386896	NaN	0	2020-02-27		Thu	22	24
24386897	NaN	0	2020-02-28		Fri	7	43
24386898	NaN	0	2020-02-28		Fri	20	10
24386899	NaN	0	2020-02-28		Fri	23	10


[24386900 rows x 19 columns]

```
[24]: # Create a fraud_df that has all of the fraudulent transactions
fraud_df = cc_df.loc[cc_df['Is Fraud?'] == 1].copy()
# Set the bins for transaction amounts that will be used in graphing
bins = bin_edges = range(0, int(fraud_df['Amount'].max()) + 100, 100)
fraud_df['Amount Bin'] = pd.cut(fraud_df['Amount'], bins=bin_edges)
```

```
[25]: # Get the number of fraud in each bin
fraud_count = fraud_df.groupby('Amount Bin', observed=True).size()
```

```
[26]: # Plot the figure
plt.figure(figsize=(10,6))
fraud_count.plot(kind='bar', color='skyblue', edgecolor='black')
plt.title('Count of Fraud by Transaction Amount')
```

```
plt.xlabel('Transaction Amount Bin $')
plt.ylabel('Number of Transactions')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```




From the above graph, we can see that the vast majority of fraudulent transactions are less than \$100. Let's take a look at locations next, we will start with State (which is also Country). Additionally, we have modified the Merchant State column to include Online information.

```
[27]: fraud_df['Merchant State'].unique()
```

```
[27]: array(['ONLINE', 'CA', 'NY', 'Italy', 'Haiti', 'Algeria', 'OH', 'MI',
        'OK', 'TX', 'NE', 'CO', 'AZ', 'IN', 'Nigeria', 'Mexico', 'Tuvalu',
        'TN', 'FL', 'MA', 'KS', 'GA', 'NJ', 'SC', 'WA', 'UT', 'MD', 'PA',
        'AL', 'Turkey', 'IL', 'Japan', 'NC', 'HI', 'IA', 'DE', 'ID', 'MN',
        'VA', 'MS', 'KY', 'WI', 'CT', 'VT', 'AR', 'NV', 'Fiji', 'NM', 'LA',
        'OR', 'MO', 'AK', 'SD', 'ND', 'ME', 'WV', 'NH', 'India', 'WY',
        'MT', 'RI'], dtype=object)
```

```
[28]: # Number of fraudulent transactions by State/Country
plt.figure(figsize=(15, 10))
sns.countplot(data=fraud_df, y='Merchant State', order=fraud_df['Merchant_
↪State'].value_counts().index)
plt.title('Number of Fraudulent Transactions by State')
plt.show()
```



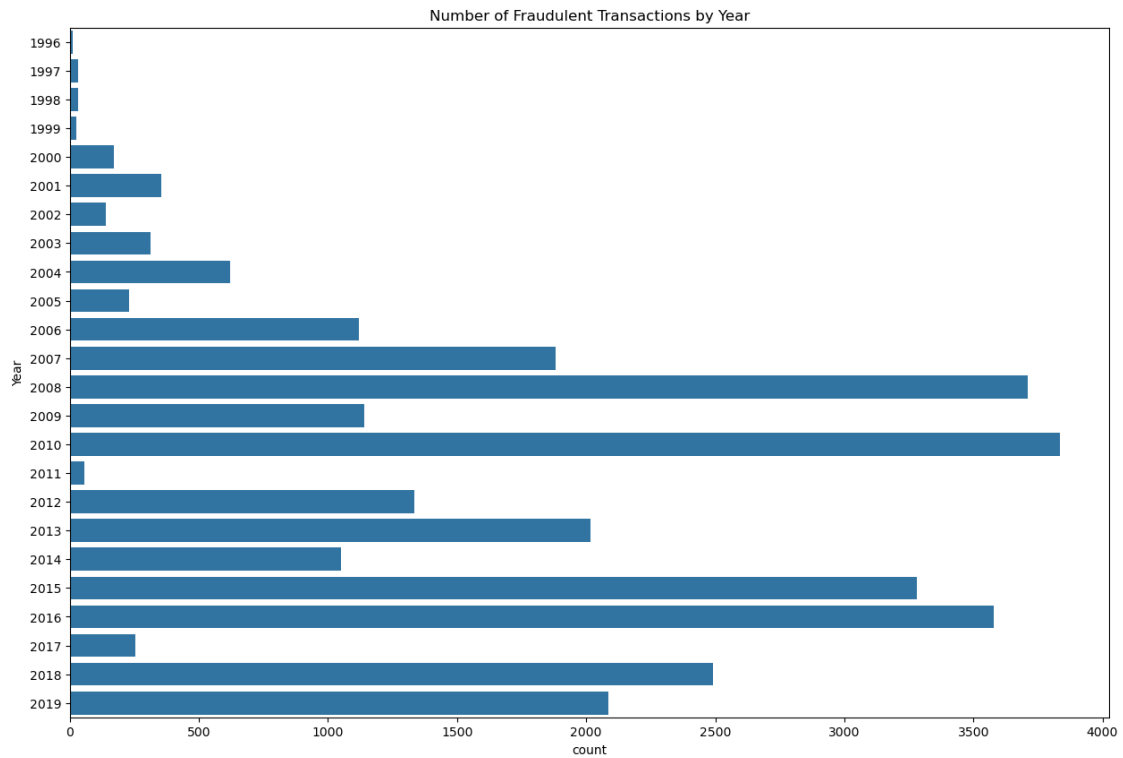
As we can see from the graph above, the majority of fraudulent transactions were made online. In second place we have Italy, and in third Ohio.

```
[29]: fraud_df['Year'].unique()
```

```
[29]: array([2015, 2016, 2008, 2019, 2010, 2006, 2018, 2013, 2017, 2001, 2014,
          2009, 2007, 2005, 2000, 2012, 2002, 1999, 2004, 2003, 2011, 1997,
          1998, 1996])
```


0.6 Date/Time Fraud Analysis

```
[30]: plt.figure(figsize=(15, 10))
      sns.countplot(data=fraud_df, y='Year')
      plt.title('Number of Fraudulent Transactions by Year')
      plt.show()
```



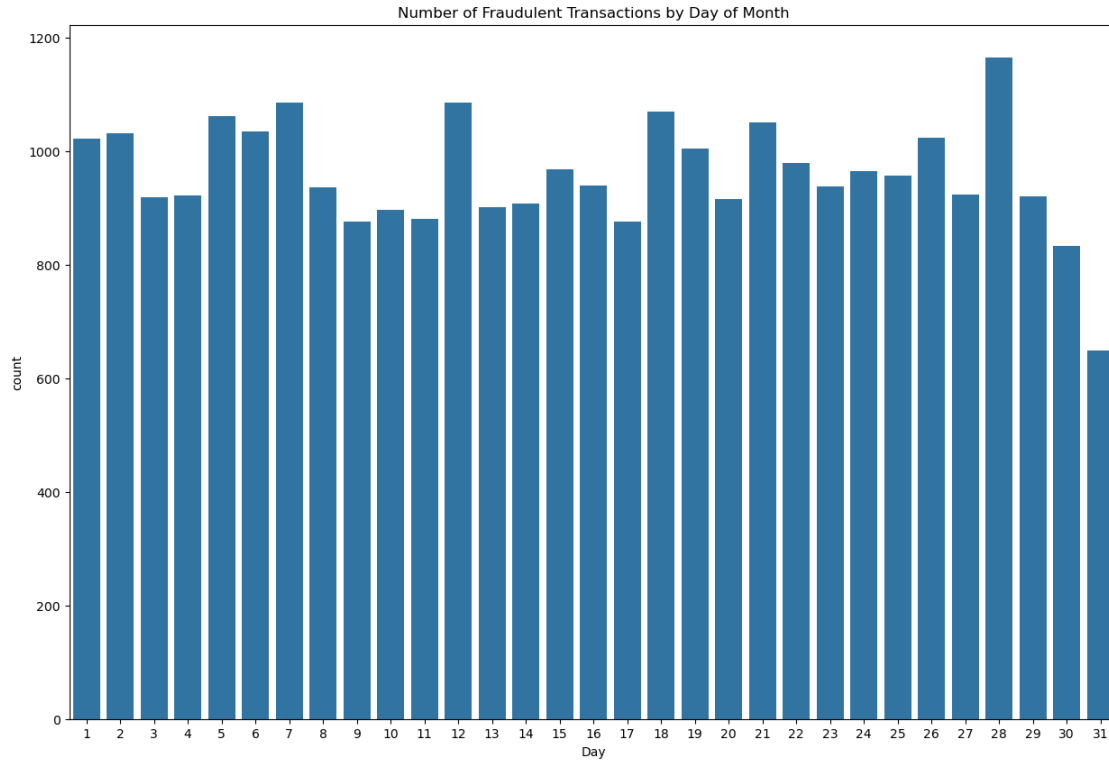
We can see that the majority of fraudulent transactions occurred in 2010, 2008, and 2016.

```
[31]: plt.figure(figsize=(15, 10))
sns.countplot(data=fraud_df, x='Month')
plt.title('Number of Fraudulent Transactions by Month')
plt.show()
```


Here, we can see that the end of the year has slightly higher rates of fraudulent transactions, but it is not a very clear trend.

```
[32]: # Graph for Day of the Month
plt.figure(figsize=(15, 10))
sns.countplot(data=fraud_df, x='Day')
plt.title('Number of Fraudulent Transactions by Day of Month')
plt.show()
```



We can see that while the 28th has the highest rate, there are not clearly defined patterns of increased fraud associated with a particular day of the month.

```
[33]: plt.figure(figsize=(15, 10))
sns.countplot(data=fraud_df, x='Day of Week', order=['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun'])
plt.title('Number of Fraudulent Transactions by Day of Week')
plt.show()
```




From the above graph, we can see that Sunday, Thursday, and Friday all see elevated rates of fraud.

```
[34]: len(fraud_df['Hour'].unique())
```


```
[34]: 24
```

```
[35]: # Let's look at fraudulent activity based on the hours of the day
plt.figure(figsize=(15, 10))
sns.countplot(data=fraud_df, x='Hour', order=range(0, 24)) # 0 to 24 inclusive
plt.title('Number of Fraudulent Transactions by Hour')
plt.xlabel('Hour of day')
plt.ylabel('Count of Fraudulent Transactions')
plt.xticks(range(0, 24), labels=range(0, 24))
plt.show()
```




From the graph above, we can see that most of fraudulent transactions occur between 10-12 in this dataset.

```
[36]: # Graph of use chip
plt.figure(figsize=(15, 10))
sns.countplot(data=fraud_df, x='Use Chip', order=fraud_df['Use Chip'].
    ↳value_counts().index)
plt.title('Number of Fraudulent Transactions by Use Chip')
plt.show()
```



```
[37]: plt.figure(figsize=(15, 20))
sns.countplot(data=fraud_df, y='MCC', order=fraud_df['MCC'].value_counts().
↪index)
plt.title('Number of Fraudulent Transactions by MCC')
plt.show()
```



0.7 Feature Selection

Looking at the above analysis, the following features are good candidates to include in our models:

Year
Amount
State
Day of Week
Hour
Use Chip
MCC

```
[87]: # Build the LSTM Model
def build_lstm_model(input_shape):
    model = Sequential()
    # Add the first LSTM layer
    model.add(LSTM(200, activation='tanh', input_shape=input_shape,
    ↪return_sequences=True))
    model.add(Dropout(0.3)) # Dropout to avoid overfitting
    # Add the second LSTM layer
    model.add(LSTM(200, activation='tanh', return_sequences=False))
    # Dense layers
    model.add(Dense(32, activation='relu'))
    model.add(Dropout(0.2))
    # Output layer
    model.add(Dense(1, activation='sigmoid'))
    model.compile(optimizer=Adam(), loss='binary_crossentropy',
    ↪metrics=['accuracy'])
    return model
```

0.8 Online Transactions Model

Since we are dealing with a heavily imbalanced dataset, let's investigate whether we can improve our model's balance by focusing exclusively on online transactions where most of the fraud occurs.

```
[120]: cc_df
```

```
[120]:
```

	User	Card	Year	Month	Day	Time	Amount	Use Chip \
0	0	0	2002	9	1	06:21	134.09	Swipe Transaction
1	0	0	2002	9	1	06:42	38.48	Swipe Transaction
2	0	0	2002	9	2	06:22	120.34	Swipe Transaction
3	0	0	2002	9	2	17:45	128.95	Swipe Transaction
4	0	0	2002	9	3	06:23	104.71	Swipe Transaction
...
24386895	1999	1	2020	2	27	22:23	-54.00	Chip Transaction
24386896	1999	1	2020	2	27	22:24	54.00	Chip Transaction
24386897	1999	1	2020	2	28	07:43	59.15	Chip Transaction
24386898	1999	1	2020	2	28	20:10	43.12	Chip Transaction
24386899	1999	1	2020	2	28	23:10	45.13	Chip Transaction

	Merchant Name	Merchant City	Merchant State	Zip	MCC	\
0	3527213246127876953	La Verne	CA	91750.0	5300	
1	-727612092139916043	Monterey Park	CA	91754.0	5411	
2	-727612092139916043	Monterey Park	CA	91754.0	5411	
3	3414527459579106770	Monterey Park	CA	91754.0	5651	
4	5817218446178736267	La Verne	CA	91750.0	5912	
...	
24386895	-5162038175624867091	Merrimack	NH	3054.0	5541	
24386896	-5162038175624867091	Merrimack	NH	3054.0	5541	
24386897	2500998799892805156	Merrimack	NH	3054.0	4121	
24386898	2500998799892805156	Merrimack	NH	3054.0	4121	
24386899	4751695835751691036	Merrimack	NH	3054.0	5814	

	Errors?	Is Fraud?	Date	Day of Week	Hour	Minute
0	NaN	0	2002-09-01	Sun	6	21
1	NaN	0	2002-09-01	Sun	6	42
2	NaN	0	2002-09-02	Mon	6	22
3	NaN	0	2002-09-02	Mon	17	45
4	NaN	0	2002-09-03	Tue	6	23
...
24386895	NaN	0	2020-02-27	Thu	22	23
24386896	NaN	0	2020-02-27	Thu	22	24
24386897	NaN	0	2020-02-28	Fri	7	43
24386898	NaN	0	2020-02-28	Fri	20	10
24386899	NaN	0	2020-02-28	Fri	23	10

[24386900 rows x 19 columns]

```
[121]: online_model_cols = ['User', 'Year', 'Day of Week', 'Hour', 'Amount', 'Use_
↳Chip', 'MCC', 'Is Fraud?']
online_df = cc_df.loc[:,online_model_cols]
online_df = online_df.loc[cc_df['Merchant State'] == 'ONLINE']
```

```
[122]: print(f"There are {cc_df.shape[0]} records in the cc_df and {online_df.
↳shape[0]} records in the online_df.")
print(f"We have eliminated {cc_df.shape[0] - online_df.shape[0]} records by_
↳using the online_df.")
```


There are 24386900 records in the cc_df and 2720821 records in the online_df.
We have eliminated 21666079 records by using the online_df.

```
[123]: online_df.shape
```

```
[123]: (2720821, 8)
```



```
[124]: # Fraudulent vs legitimate transactions
plt.figure(figsize=(15, 10))
sns.countplot(data=online_df, x='Is Fraud?')
plt.title('Fraudulent vs Legitimate Online Transactions')
plt.show()
```



0.9 Online Model Create - Random Forest

```
[180]: online_rf_df = online_df.copy()
online_rf_df
```

```
[180]:
```

	User	Year	Day	of Week	Hour	Amount	Use Chip	MCC	\
11	0	2002		Thu	20	53.91	Online Transaction	4900	
24	0	2002		Mon	20	144.90	Online Transaction	4899	
85	0	2002		Mon	6	127.32	Online Transaction	5311	
99	0	2002		Sun	6	139.39	Online Transaction	5311	
106	0	2002		Wed	8	53.09	Online Transaction	5193	
...	
24386877	1999	2020		Mon	20	55.79	Online Transaction	4121	
24386879	1999	2020		Tue	7	43.08	Online Transaction	4121	
24386880	1999	2020		Tue	7	43.76	Online Transaction	4121	
24386884	1999	2020		Wed	7	45.18	Online Transaction	4121	

24386889	1999	2020	Thu	7	47.18	Online Transaction	4121
----------	------	------	-----	---	-------	--------------------	------

	Is Fraud?
11	0
24	0
85	0
99	0
106	0
...	...
24386877	0
24386879	0
24386880	0
24386884	0
24386889	0

[2720821 rows x 8 columns]

```
[181]: # Target encode the user features
user_target_mean = online_rf_df.groupby('User')['Is Fraud?'].mean()
online_rf_df['User_encoded'] = online_rf_df['User'].map(user_target_mean)
online_rf_df = online_rf_df.drop(columns=['User'])
# Convert columns to the appropriate dtype
online_rf_df['MCC'] = online_rf_df['MCC'].astype('category')

# Encode the Object columns accordingly
online_rf_df = pd.get_dummies(online_rf_df, columns=['Day of Week', 'Use_
↳Chip'], prefix="Day")

# Create a binary encoder
rf_binary_encoder = ce.BinaryEncoder(cols=['MCC'])
online_rf_df = rf_binary_encoder.fit_transform(online_rf_df)

# Sine-Cosine Encoding for Hour and Year
online_rf_df['Hour_Sin'] = np.sin(2 * np.pi * online_rf_df['Hour'] / 24)
online_rf_df['Hour_Cos'] = np.cos(2 * np.pi * online_rf_df['Hour'] / 24)
online_rf_df.drop(columns='Hour', inplace=True)

rf_target_column = "Is Fraud?"

online_rf_df
```

[181]:	Year	Amount	MCC_0	MCC_1	MCC_2	MCC_3	MCC_4	MCC_5	MCC_6	\
11	2002	53.91	0	0	0	0	0	0	1	
24	2002	144.90	0	0	0	0	0	1	0	
85	2002	127.32	0	0	0	0	0	1	1	
99	2002	139.39	0	0	0	0	0	1	1	
106	2002	53.09	0	0	0	0	1	0	0	

...
24386877	2020	55.79	0	0	1	0	1	0	0	0
24386879	2020	43.08	0	0	1	0	1	0	0	0
24386880	2020	43.76	0	0	1	0	1	0	0	0
24386884	2020	45.18	0	0	1	0	1	0	0	0
24386889	2020	47.18	0	0	1	0	1	0	0	0

	Is_Fraud?	...	Day_Mon	Day_Sat	Day_Sun	Day_Thu	Day_Tue	\
11	0	...	False	False	False	True	False	
24	0	...	True	False	False	False	False	
85	0	...	True	False	False	False	False	
99	0	...	False	False	True	False	False	
106	0	...	False	False	False	False	False	

...
24386877	0	...	True	False	False	False	False	False
24386879	0	...	False	False	False	False	False	True
24386880	0	...	False	False	False	False	False	True
24386884	0	...	False	False	False	False	False	False
24386889	0	...	False	False	False	True	False	False

	Day_Wed	Day_Chip	Transaction	Day_Online	Transaction	Hour_Sin	\
11	False		False		True	-0.866025	
24	False		False		True	-0.866025	
85	False		False		True	1.000000	
99	False		False		True	1.000000	
106	True		False		True	0.866025	

...
24386877	False		False		True	-0.866025	
24386879	False		False		True	0.965926	
24386880	False		False		True	0.965926	
24386884	True		False		True	0.965926	
24386889	False		False		True	0.965926	

	Hour_Cos
11	5.000000e-01
24	5.000000e-01
85	6.123234e-17
99	6.123234e-17
106	-5.000000e-01

...	...
24386877	5.000000e-01
24386879	-2.588190e-01
24386880	-2.588190e-01
24386884	-2.588190e-01
24386889	-2.588190e-01

[2720821 rows x 22 columns]

```
[182]: # Split the features from the target variable
rf_X = online_rf_df.drop(columns=[rf_target_column])
rf_y = online_rf_df[rf_target_column]

# Train Test Split
rf_X_train, rf_X_test, rf_y_train, rf_y_test = train_test_split(rf_X, rf_y,
    ↪test_size=0.2, random_state=42, stratify=rf_y)
```

```
[183]: # Create RF Model
rf_model = RandomForestClassifier(n_estimators=100, random_state=42,
    ↪class_weight='balanced', n_jobs=-1)

# Fit the model
rf_model.fit(rf_X_train, rf_y_train)
```

```
[183]: RandomForestClassifier(class_weight='balanced', n_jobs=-1, random_state=42)
```

```
[184]: # Evaluate the model
rf_y_pred = rf_model.predict(rf_X_test)
rf_y_pred_proba = rf_model.predict_proba(rf_X_test)[: , 1]
```

```
[191]: rf_cm = confusion_matrix(rf_y_test, rf_y_pred)
print(f"Classification Report:\n{classification_report(rf_y_test, rf_y_pred)}")
print(f"Confusion Matrix:\n{rf_cm}")
print(f"ROC AUC Score:\n {roc_auc_score(rf_y_test, rf_y_pred_proba)}")
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	540495
1	0.97	0.63	0.76	3670
accuracy			1.00	544165
macro avg	0.99	0.81	0.88	544165
weighted avg	1.00	1.00	1.00	544165

Confusion Matrix:


```
[[540430    65]
 [ 1366   2304]]
```

ROC AUC Score:

```
0.9812301913275431
```

```
[196]: # Plot the confusion matrix
plt.figure(figsize=(7,5))
annotations = [f'TN: {rf_cm[0,0]}', f'FP: {online_cm[0,1]}'], [f'FN:
    ↪{rf_cm[1,0]}', f'TP: {rf_cm[1,1]}']
hm = sns.heatmap(rf_cm, annot=True, fmt='d', cmap='Reds',
    ↪xticklabels=['Non-Fraud', 'Fraud'], yticklabels=['Non-Fraud', 'Fraud'])
```


```
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix (Random Forest)')
hm.text(0.1, 0.2, 'True Negative', color='white')
hm.text(1.1, 0.2, 'False Positive')
hm.text(0.1, 1.2, 'False Negative')
hm.text(1.1, 1.2, 'True Positive')
plt.show()
```



```
[192]: # Get the ROC Curve
rf_fpr, rf_tpr, rf_thresholds = roc_curve(rf_y_test, rf_y_pred_proba)
rf_roc_auc = auc(rf_fpr, rf_tpr)
```

```
[195]: # Plot the ROC curve
plt.figure()
plt.plot(rf_fpr, rf_tpr, label=f'ROC Curve (area = {online_roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--') # Random guess line
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (Random Forest)')
```

```
plt.legend(loc='lower right')
plt.show()
```



0.10 Online Model Creation - LSTM

```
[198]: # Create the dataset to be used by the models
online_model_df = online_df.copy()
```

0.11 Encode the data

```
[126]: # Target encode the user features
user_target_mean = online_model_df.groupby('User')['Is Fraud?'].mean()
online_model_df['User_encoded'] = online_model_df['User'].map(user_target_mean)
online_model_df = online_model_df.drop(columns=['User'])
# Convert columns to the appropriate dtype
online_model_df['MCC'] = online_model_df['MCC'].astype('category')

# Encode the Object columns accordingly
online_model_df = pd.get_dummies(online_model_df, columns=['Day of Week', 'Use_
↳ Chip'], prefix="Day")
```

```

# Create a binary encoder
online_binary_encoder = ce.BinaryEncoder(cols=['MCC'])
online_model_df = online_binary_encoder.fit_transform(online_model_df)

# Sine-Cosine Encoding for Hour and Year
online_model_df['Hour_Sin'] = np.sin(2 * np.pi * online_model_df['Hour'] / 24)
online_model_df['Hour_Cos'] = np.cos(2 * np.pi * online_model_df['Hour'] / 24)
online_model_df.drop(columns='Hour', inplace=True)

online_target_column = "Is Fraud?"

online_model_df

```

```

[126]:
      Year  Amount  MCC_0  MCC_1  MCC_2  MCC_3  MCC_4  MCC_5  MCC_6  \
11      2002   53.91      0      0      0      0      0      0      1
24      2002  144.90      0      0      0      0      0      1      0
85      2002  127.32      0      0      0      0      0      1      1
99      2002  139.39      0      0      0      0      0      1      1
106     2002   53.09      0      0      0      0      1      0      0
...
24386877  2020   55.79      0      0      1      0      1      0      0
24386879  2020   43.08      0      0      1      0      1      0      0
24386880  2020   43.76      0      0      1      0      1      0      0
24386884  2020   45.18      0      0      1      0      1      0      0
24386889  2020   47.18      0      0      1      0      1      0      0

      Is Fraud?  ...  Day_Mon  Day_Sat  Day_Sun  Day_Thu  Day_Tue  \
11              0  ...   False   False   False    True   False
24              0  ...    True   False   False   False   False
85              0  ...    True   False   False   False   False
99              0  ...   False   False    True   False   False
106             0  ...   False   False   False   False   False
...
24386877         0  ...    True   False   False   False   False
24386879         0  ...   False   False   False   False    True
24386880         0  ...   False   False   False   False    True
24386884         0  ...   False   False   False   False   False
24386889         0  ...   False   False   False    True   False

      Day_Wed  Day_Chip  Transaction  Day_Online  Transaction  Hour_Sin  \
11         False              False              True -0.866025
24         False              False              True -0.866025
85         False              False              True  1.000000
99         False              False              True  1.000000
106        True              False              True  0.866025
...

```

24386877	False	False	True	-0.866025
24386879	False	False	True	0.965926
24386880	False	False	True	0.965926
24386884	True	False	True	0.965926
24386889	False	False	True	0.965926

	Hour_Cos
11	5.000000e-01
24	5.000000e-01
85	6.123234e-17
99	6.123234e-17
106	-5.000000e-01
...	...
24386877	5.000000e-01
24386879	-2.588190e-01
24386880	-2.588190e-01
24386884	-2.588190e-01
24386889	-2.588190e-01

[2720821 rows x 22 columns]

```
[127]: # Use compute_class_weights to handle class imbalance
online_y_train = online_df['Is Fraud?']

# Compute class weights
class_weights = compute_class_weight(class_weight='balanced', classes=np.
    ↳unique(online_y_train), y=online_y_train)

# Get Class Weights
class_weights_dict = {i: weight for i, weight in enumerate(class_weights)}
class_weights_dict
```

```
[127]: {0: 0.5033948547848044, 1: 74.14085236252657}
```

```
[128]: # Split the features from the target variable
online_X = online_model_df.drop(columns=['Is Fraud?']) # Features
online_y = online_model_df['Is Fraud?'] # Target

# Create the train and test sets
online_X_train, online_X_test, online_y_train, online_y_test =
    ↳train_test_split(online_X, online_y, test_size=0.2, random_state=42)

# Scale the data for the Random Forest Classifier and LSTM
online_scaler = StandardScaler()
online_X_train_scaled = online_scaler.fit_transform(online_X_train)
online_X_test_scaled = online_scaler.transform(online_X_test)
```



```
[129]: # Make sure that my GPU is available
print(f"Num GPUs available: {len(tf.config.list_physical_devices('GPU'))}")

# Reshape the data for LSTM
online_X_train_lstm = np.reshape(online_X_train_scaled, (online_X_train_scaled.
↳shape[0], 1, online_X_train_scaled.shape[1]))
online_X_test_lstm = np.reshape(online_X_test_scaled, (online_X_test_scaled.
↳shape[0], 1, online_X_test_scaled.shape[1]))
```

Num GPUs available: 1

```
[130]: # Build and train the model
online_input_shape = (online_X_train_lstm.shape[1], online_X_train_lstm.
↳shape[2])
online_lstm_model = build_lstm_model(online_input_shape)
```

```
[131]: online_lstm_model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
lstm_2 (LSTM)	(None, 1, 200)	177600
dropout_2 (Dropout)	(None, 1, 200)	0
lstm_3 (LSTM)	(None, 200)	320800
dense_2 (Dense)	(None, 32)	6432
dropout_3 (Dropout)	(None, 32)	0
dense_3 (Dense)	(None, 1)	33

```
=====
Total params: 504865 (1.93 MB)
Trainable params: 504865 (1.93 MB)
Non-trainable params: 0 (0.00 Byte)
=====
```

```
[132]: # Time the training process
online_start_time = time.time()
online_history = online_lstm_model.fit(online_X_train_lstm, online_y_train,
↳epochs=10, batch_size=64, validation_split=0.2,
↳class_weight=class_weights_dict)
online_end_time = time.time()

print(f"Training time: {online_end_time - online_start_time:.2f} seconds")
```

```

Epoch 1/10
27209/27209 [=====] - 308s 11ms/step - loss: 0.2595 -
accuracy: 0.8695 - val_loss: 0.1754 - val_accuracy: 0.9113
Epoch 2/10
27209/27209 [=====] - 314s 12ms/step - loss: 0.1802 -
accuracy: 0.9113 - val_loss: 0.1458 - val_accuracy: 0.9260
Epoch 3/10
27209/27209 [=====] - 309s 11ms/step - loss: 0.1726 -
accuracy: 0.9169 - val_loss: 0.1390 - val_accuracy: 0.9332
Epoch 4/10
27209/27209 [=====] - 302s 11ms/step - loss: 0.1695 -
accuracy: 0.9169 - val_loss: 0.1596 - val_accuracy: 0.9135
Epoch 5/10
27209/27209 [=====] - 309s 11ms/step - loss: 0.1681 -
accuracy: 0.9176 - val_loss: 0.1545 - val_accuracy: 0.9224
Epoch 6/10
27209/27209 [=====] - 306s 11ms/step - loss: 0.1666 -
accuracy: 0.9180 - val_loss: 0.1568 - val_accuracy: 0.9269
Epoch 7/10
27209/27209 [=====] - 296s 11ms/step - loss: 0.1799 -
accuracy: 0.9174 - val_loss: 0.1485 - val_accuracy: 0.9274
Epoch 8/10
27209/27209 [=====] - 288s 11ms/step - loss: 0.1904 -
accuracy: 0.9166 - val_loss: 0.1340 - val_accuracy: 0.9292
Epoch 9/10
27209/27209 [=====] - 281s 10ms/step - loss: 0.1973 -
accuracy: 0.9160 - val_loss: 0.1347 - val_accuracy: 0.9296
Epoch 10/10
27209/27209 [=====] - 283s 10ms/step - loss: 0.2345 -
accuracy: 0.9128 - val_loss: 0.2338 - val_accuracy: 0.8977
Training time: 2997.02 seconds

```

```

[133]: online_loss, online_accuracy = online_lstm_model.evaluate(online_X_test_lstm,
↪online_y_test, batch_size=64, verbose=1)

```

```

8503/8503 [=====] - 39s 5ms/step - loss: 0.2332 -
accuracy: 0.8981

```

```

[134]: print(f"Test loss: {online_loss:.4f}")
print(f"Test Accuracy: {online_accuracy:.4f}")

```

```

Test loss: 0.2332
Test Accuracy: 0.8981

```

```

[135]: online_lstm_pred = online_lstm_model.predict(online_X_test_lstm)

```

```

17006/17006 [=====] - 46s 3ms/step

```

```
[136]: # Convert probabilities to binary predictions (0 or 1)
threshold = 0.5
online_binary_predictions = (online_lstm_pred > threshold).astype(int)
```

```
[137]: # Classification Report
print(classification_report(online_y_test, online_binary_predictions))


# Confusion Matrix
print(f"Confusion Matrix:\n{confusion_matrix(online_y_test,
↪online_binary_predictions)}")
```

	precision	recall	f1-score	support
0	1.00	0.90	0.95	540459
1	0.06	0.97	0.11	3706
accuracy			0.90	544165
macro avg	0.53	0.93	0.53	544165
weighted avg	0.99	0.90	0.94	544165

Confusion Matrix:


```
[[485110  55349]
 [   116   3590]]
```

```
[197]: online_cm = confusion_matrix(online_y_test, online_binary_predictions)
# Plot the confusion matrix
plt.figure(figsize=(7,5))
annotations = [f'TN: {online_cm[0,0]}', f'FP: {online_cm[0,1]}'], [f'FN:
↪{online_cm[1,0]}', f'TP: {online_cm[1,1]}']
hm = sns.heatmap(online_cm, annot=True, fmt='d', cmap='Reds',
↪xticklabels=['Non-Fraud', 'Fraud'], yticklabels=['Non-Fraud', 'Fraud'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix (LSTM)')
hm.text(0.1, 0.2, 'True Negative', color='white')
hm.text(1.1, 0.2, 'False Positive')
hm.text(0.1, 1.2, 'False Negative')
hm.text(1.1, 1.2, 'True Positive')
plt.show()
```




```
[138]: import matplotlib.pyplot as plt

# Plot loss
plt.plot(online_history.history['loss'], label='Training Loss')
plt.plot(online_history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

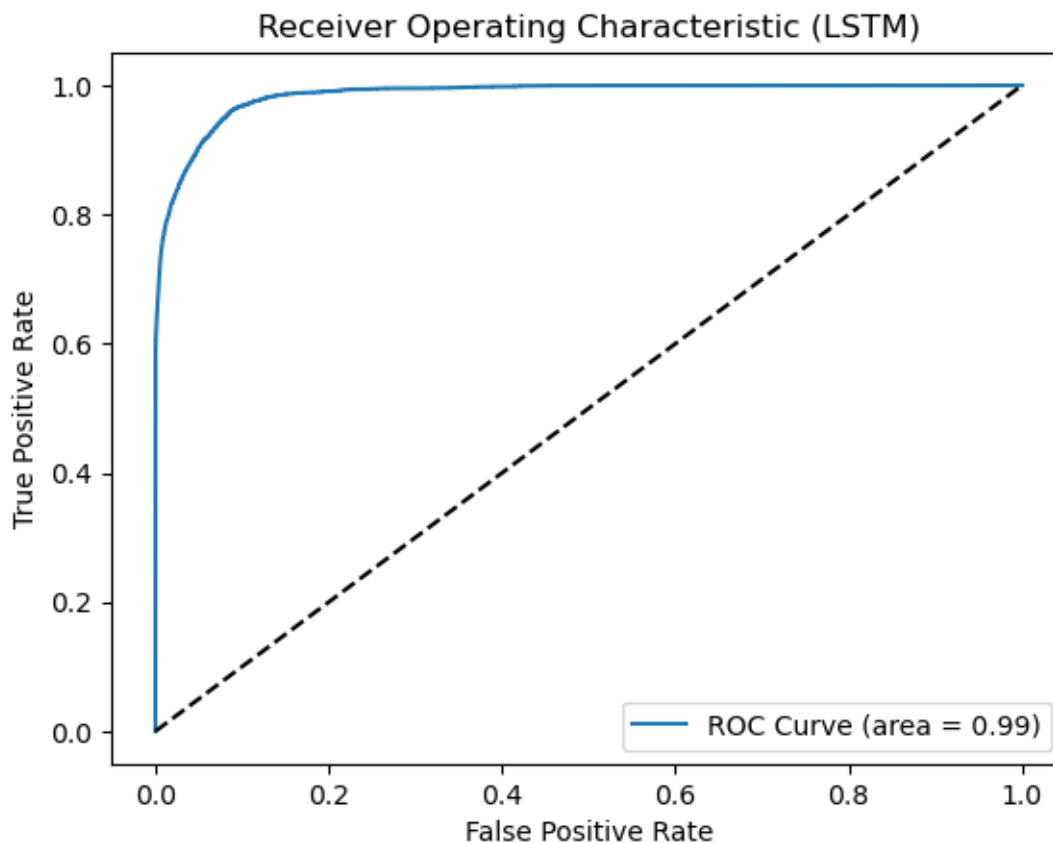


```
[139]: # Plot accuracy
plt.plot(online_history.history['accuracy'], label='Training Accuracy')
plt.plot(online_history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



```
[194]: # Get the ROC curve
online_fpr, online_tpr, online_thresholds = roc_curve(online_y_test,
    ↪online_lstm_pred)
online_roc_auc = auc(online_fpr, online_tpr)

# Plot the ROC curve
plt.figure()
plt.plot(online_fpr, online_tpr, label=f'ROC Curve (area = {online_roc_auc:.
    ↪2f})')
plt.plot([0, 1], [0, 1], 'k--') # Random guess line
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (LSTM)')
plt.legend(loc='lower right')
plt.show()
```



```
[141]: # Save the model
online_lstm_model.save('online_lstm_fraud_model2_user.keras')

# Loading the model
# from tensorflow.keras.models import load_model
# model = load_model('lstm_fraud_model.h5')
```

0.12 Experimental: Random Undersampling

```
[42]: cc_df
```

```
[42]:
```

	User	Card	Year	Month	Day	Time	Amount	Use Chip \
0	0	0	2002	9	1	06:21	134.09	Swipe Transaction
1	0	0	2002	9	1	06:42	38.48	Swipe Transaction
2	0	0	2002	9	2	06:22	120.34	Swipe Transaction
3	0	0	2002	9	2	17:45	128.95	Swipe Transaction
4	0	0	2002	9	3	06:23	104.71	Swipe Transaction
...
24386895	1999	1	2020	2	27	22:23	-54.00	Chip Transaction
24386896	1999	1	2020	2	27	22:24	54.00	Chip Transaction

24386897	1999	1	2020	2	28	07:43	59.15	Chip Transaction
24386898	1999	1	2020	2	28	20:10	43.12	Chip Transaction
24386899	1999	1	2020	2	28	23:10	45.13	Chip Transaction

	Merchant Name	Merchant City	Merchant State	Zip	MCC	\
0	3527213246127876953	La Verne	CA	91750.0	5300	
1	-727612092139916043	Monterey Park	CA	91754.0	5411	
2	-727612092139916043	Monterey Park	CA	91754.0	5411	
3	3414527459579106770	Monterey Park	CA	91754.0	5651	
4	5817218446178736267	La Verne	CA	91750.0	5912	
...	
24386895	-5162038175624867091	Merrimack	NH	3054.0	5541	
24386896	-5162038175624867091	Merrimack	NH	3054.0	5541	
24386897	2500998799892805156	Merrimack	NH	3054.0	4121	
24386898	2500998799892805156	Merrimack	NH	3054.0	4121	
24386899	4751695835751691036	Merrimack	NH	3054.0	5814	

	Errors?	Is Fraud?	Date	Day of Week	Hour	Minute
0	NaN	0	2002-09-01	Sun	6	21
1	NaN	0	2002-09-01	Sun	6	42
2	NaN	0	2002-09-02	Mon	6	22
3	NaN	0	2002-09-02	Mon	17	45
4	NaN	0	2002-09-03	Tue	6	23
...
24386895	NaN	0	2020-02-27	Thu	22	23
24386896	NaN	0	2020-02-27	Thu	22	24
24386897	NaN	0	2020-02-28	Fri	7	43
24386898	NaN	0	2020-02-28	Fri	20	10
24386899	NaN	0	2020-02-28	Fri	23	10

[24386900 rows x 19 columns]

[43]: user_df

[43]:

	Person	Current Age	Retirement Age	Birth Year	Birth Month	\
0	Hazel Robinson	53	66	1966	11	
1	Sasha Sadr	53	68	1966	12	
2	Saanvi Lee	81	67	1938	11	
3	Everlee Clark	63	63	1957	1	
4	Kyle Peterson	43	70	1976	9	
...	
1995	Jose Faraday	32	70	1987	7	
1996	Ximena Richardson	62	65	1957	11	
1997	Annika Russell	47	67	1973	1	
1998	Juelz Roman	66	60	1954	2	
1999	Kenia Harris	21	60	1998	11	

	Gender	Address	Apartment	City	State	\
0	Female	462 Rose Lane	NaN	La Verne	CA	
1	Female	3606 Federal Boulevard	NaN	Little Neck	NY	
2	Female	766 Third Drive	NaN	West Covina	CA	
3	Female	3 Madison Street	NaN	New York	NY	
4	Male	9620 Valley Stream Drive	NaN	San Francisco	CA	
...	
1995	Male	6577 Lexington Lane	9.0	Freeport	NY	
1996	Female	2 Elm Drive	955.0	Independence	KY	
1997	Female	276 Fifth Boulevard	NaN	Elizabeth	NJ	
1998	Male	259 Valley Boulevard	NaN	Camp Hill	PA	
1999	Female	472 Ocean View Street	NaN	Merrimack	NH	

	Zipcode	Latitude	Longitude	Per Capita Income	- Zipcode	\
0	91750	34.15	-117.76		\$29278	
1	11363	40.76	-73.74		\$37891	
2	91792	34.02	-117.89		\$22681	
3	10069	40.71	-73.99		\$163145	
4	94117	37.76	-122.44		\$53797	
...		
1995	11520	40.65	-73.58		\$23550	
1996	41051	38.95	-84.54		\$24218	
1997	7201	40.66	-74.19		\$15175	
1998	17011	40.24	-76.92		\$25336	
1999	3054	42.86	-71.48		\$32325	

	Yearly Income - Person	Total Debt	FICO Score	Num Credit Cards
0	\$59696	\$127613	787	5
1	\$77254	\$191349	701	5
2	\$33483	\$196	698	5
3	\$249925	\$202328	722	4
4	\$109687	\$183855	675	1
...
1995	\$48010	\$87837	703	3
1996	\$49378	\$104480	740	4
1997	\$30942	\$71066	779	3
1998	\$54654	\$27241	618	1
1999	\$65909	\$181261	673	2

[2000 rows x 18 columns]

```
[172]: select_user_cols = ['Current Age', 'Gender', 'State', 'Yearly Income - Person', 'FICO Score', 'Num Credit Cards']
```

```
[173]: user_df.columns
```

```
[173]: Index(['Person', 'Current Age', 'Retirement Age', 'Birth Year', 'Birth Month',
          'Gender', 'Address', 'Apartment', 'City', 'State', 'Zipcode',
          'Latitude', 'Longitude', 'Per Capita Income - Zipcode',
          'Yearly Income - Person', 'Total Debt', 'FICO Score',
          'Num Credit Cards'],
          dtype='object')
```

```
[175]: user_df.loc[:,select_user_cols]
```

```
[175]:
```

	Current Age	Gender	State	Yearly Income - Person	FICO Score \
0	53	Female	CA	\$59696	787
1	53	Female	NY	\$77254	701
2	81	Female	CA	\$33483	698
3	63	Female	NY	\$249925	722
4	43	Male	CA	\$109687	675
...
1995	32	Male	NY	\$48010	703
1996	62	Female	KY	\$49378	740
1997	47	Female	NJ	\$30942	779
1998	66	Male	PA	\$54654	618
1999	21	Female	NH	\$65909	673

	Num Credit Cards
0	5
1	5
2	5
3	4
4	1
...	...
1995	3
1996	4
1997	3
1998	1
1999	2

[2000 rows x 6 columns]

```
[178]: filter_user = user_df.copy()
       filter_user = user_df.loc[:, select_user_cols]
```

```
[179]: filter_user
```

```
[179]:
```

	Current Age	Gender	State	Yearly Income - Person	FICO Score \
0	53	Female	CA	\$59696	787
1	53	Female	NY	\$77254	701
2	81	Female	CA	\$33483	698
3	63	Female	NY	\$249925	722

4	43	Male	CA	\$109687	675
...
1995	32	Male	NY	\$48010	703
1996	62	Female	KY	\$49378	740
1997	47	Female	NJ	\$30942	779
1998	66	Male	PA	\$54654	618
1999	21	Female	NH	\$65909	673

	Num Credit Cards
0	5
1	5
2	5
3	4
4	1
...	...
1995	3
1996	4
1997	3
1998	1
1999	2

[2000 rows x 6 columns]